

Markov Decision Process Value Iteration

Anna Helena Reali Costa

Lecture slides partially extracted from:

<http://aima.eecs.berkeley.edu/instructors.html>

<http://ocw.mit.edu/> (course 6.825)

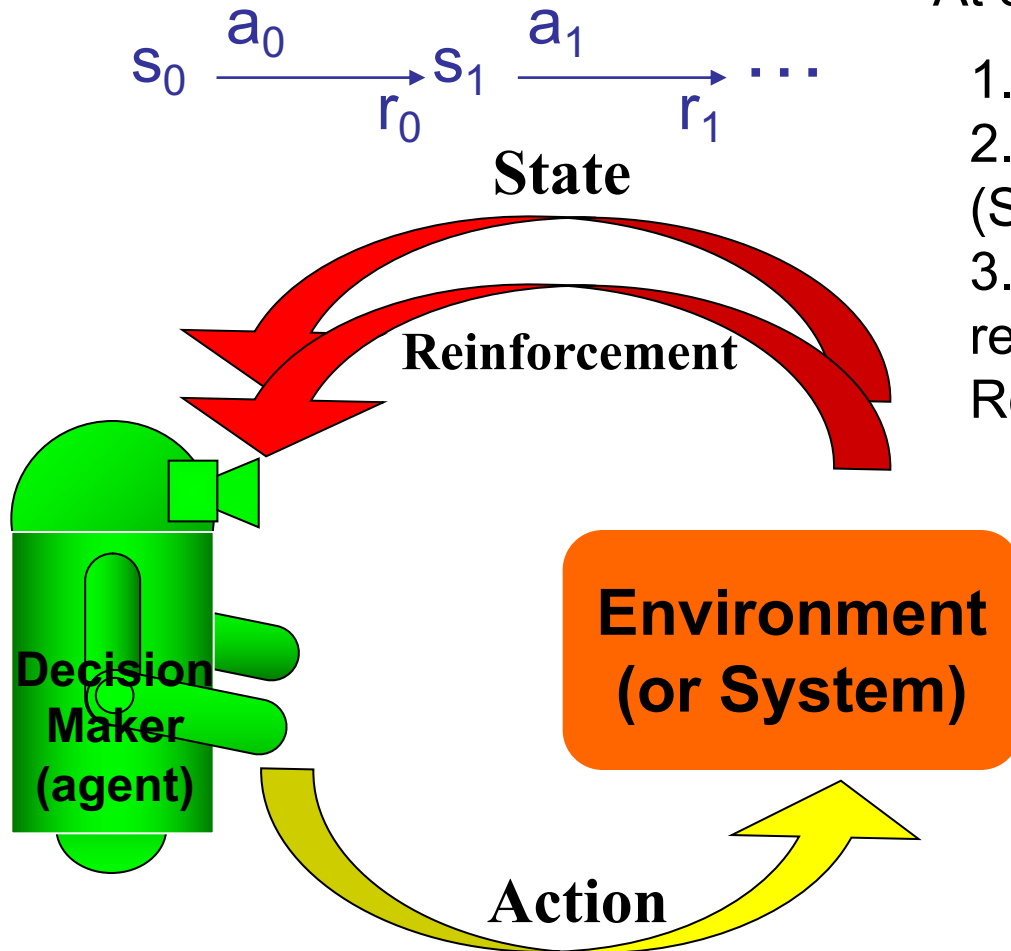
<http://www.laas.fr/planning/>

Chapter 13: Machine Learning, Tom Mitchell

Characteristics of the problem

- The agent has a set of **sensors** to observe the *state* of its environment
- The agent has a set of *actions* it can perform to alter this state
- The agent perceives a **reward** (or penalty) to indicate the desirability of the resulting state
- ❖ The task of the agent is to learn from this indirect, delayed reward, to choose **sequences of actions** that produce the greatest cumulative reward
 - ➔ it is a *sequential decision problem*

Sequential decision problem



At each time step the decision maker:


1. Observes the state of the system;
 2. Chooses an action and applies it;
(System evolves to a new state)
 3. Observes an immediate reinforcement (reward or penalty);
- Repeat 1 – 3

Sequential decision problem

$$s_0 \xrightarrow[a_0]{r_0} s_1 \xrightarrow[a_1]{r_1} \dots$$

State

Reinforcement



Decision
Maker
(agent)

Environment
(or System)

Action

At each time step the decision maker:

1. Observes the state of the system;
 2. Chooses an action and applies it;
(System evolves to a new state)
 3. Observes an immediate reinforcement;
- Repeat 1 – 3

This assumes discrete time.

Decisions are made at points of time referred to as ***decision epochs***.


The set of decision epochs can be **finite or infinite**:

$$T = \{0, 1, 2, \dots, N\}, N \leq \infty.$$

FAZER: Testinho 1 e 2

- 4×3 discrete fully-observed environment
- Actions: Up, Down, Left and Right
- Initial state: (1,1)
- Sequence [U, U, R, R, R]:
 - (i) goes up around the barrier and reaches the goal state (4,3) with probability
 - (ii) there is also a chance of accidentally reaching the goal by going (1,1) → (2,1) → (3,1) → (3,2) → (3,3) → (4,3) with probability

			+1 <i>goal</i>
			-1
<i>start</i>			



Utility function

- **Utility function** for the agent depends on a sequence of states (environment *history*)
- In each state s , the agent receives a **reinforcement** $r(s)$, which may be positive or negative, but must be **bounded**.
- **Utility = sum of the rewards received**

FAZER: testinho 3


❓ **Utility function:** sum of the rewards received

❓ Here: $\mathbf{r(s) = -0.04} \quad \forall s$ except
 $r(4,3) = +1$ and $r(4,2) = -1$

➔ reach goal after **10 steps** (avoiding (4,2)):
 utility =

-0,04	-0,04	-0,04	+1 <i>goal</i>
-0,04		-0,04	-1
-0,04	-0,04	-0,04	-0,04
<i>start</i>			

Transition model:



Ex: (1,1),U,(1,2),D,(1,1),U,(1,2),U,(1,3),L,(1,3),R,(2,3),R,(3,3),L,(2,3),R,(3,3),R,(4,3)

Modeling sequential decision problems as Markov Decision Processes


MDP – Model Formulation

An MDP is defined as $\langle S, A, p, r \rangle$:

- S is the set of possible system states (arbitrary finite set);
- A is the set of allowable actions (arbitrary finite set);
- $p: S \times A \times S \rightarrow [0,1]$ is the **transition probability function**;
- $r: S \times A \rightarrow \mathbb{R}$ is the **reinforcement function**;

MDP

- The set **A** of allowable actions:
 - ◆ $A = \cup_{s \in S} A_s$ where A_s is the set of allowable actions in state $s \in S$
 - ◆ Or we might restrict the model: $A = A_s$ for all $s \in S$
- The transition probability function **p**:
 - ◆ $p(j \mid s, a)$ --- or $p(s, a, j)$ --- denotes the probability that the system is in state $j \in S$ at time $t+1$, when the decision maker performs action $a \in A_s$ in state $s \in S$ at time t .
 - ◆ $\sum_{j \in S} p(j \mid s, a) = 1$



MDP

- The reinforcement function \mathbf{r} :
 - ◆ $r(s,a)$, denotes the value of the reward (or reinforcement or cost) received when performing $a \in A_s$ in $s \in S$ at time t .
 - ◆ When positive, $r(s,a)$ is an *income*, and when negative it is a *cost*.
 - ◆ Can be:
 - $r(s)$;
 - $r(s,a)$;
 - $r(s,a,j)$ with $r(s,a) = \sum_{j \in S} r(s,a,j) p(j \mid s, a)$

Why “Markov”?

- ❑ The qualifier *Markov* is used because the transition probability function \mathbf{p} and the reinforcement function \mathbf{r} depend on the past through the current state of the system and the action selected by the decision maker in that state.

Notation: $X_{a:b} = X_a, X_{a+1}, \dots, X_{b-1}, X_b$

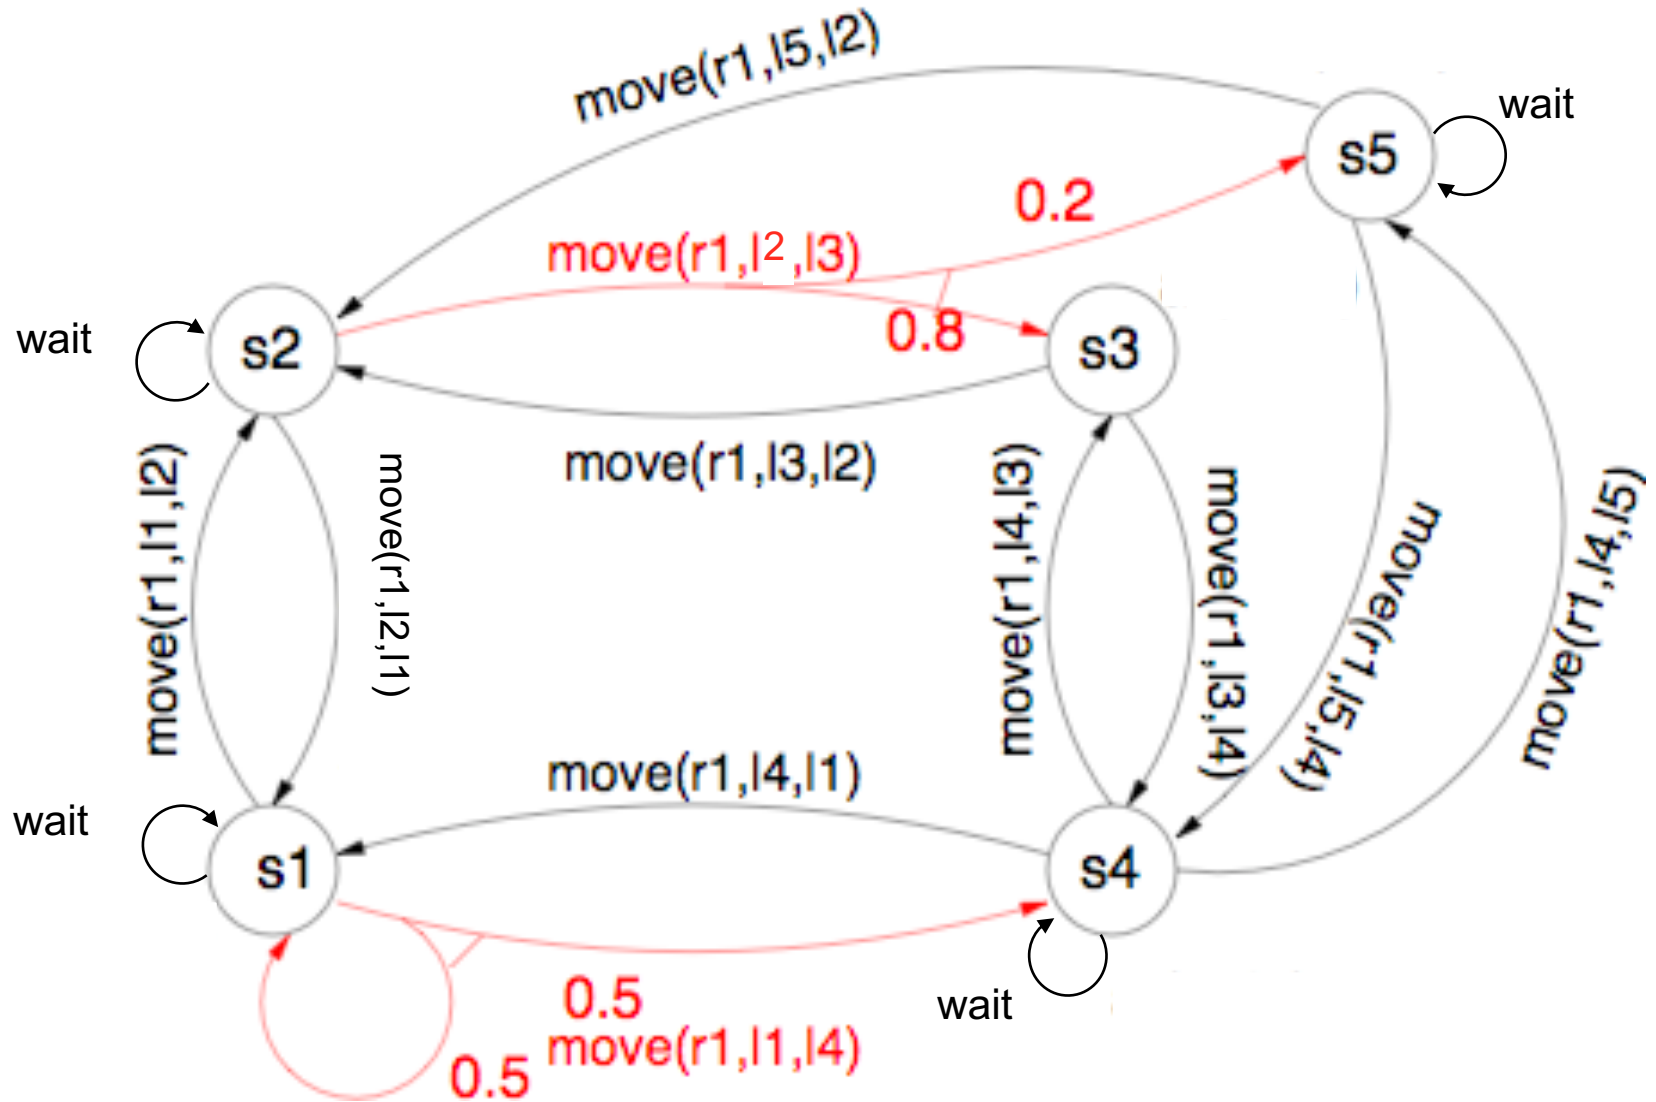
Markov assumption: X_t depends on bounded subset of $X_{0:t-1}$

- ❑ First-order Markov process: $\mathbf{p}(X_t|X_{0:t-1}) = \mathbf{p}(X_t|X_{t-1})$

Example of an MDP

- $A = \{\text{move}(r1,l1,l2), \text{move}(r1,l2,l1), \text{move}(r1,l4,l1), \text{move}(r1,l1,l4), \text{move}(r1,l3,l2), \text{move}(r1,l2,l3), \text{move}(r1,l5,l2), \text{move}(r1,l4,l3), \text{move}(r1,l3,l4), \text{move}(r1,l5,l4), \text{move}(r1,l4,l5), \text{wait}\}$
- $S = \{s1, s2, s3, s4, s5\}$
- $p(s1, \text{move}(r1,l1,l4), s4) = 0.5; p(s1, \text{move}(r1,l1,l4), s1) = 0.5;$
 $p(s2, \text{move}(r1,l2,l3), s3) = 0.8; p(s2, \text{move}(r1,l2,l3), s5) = 0.2;$
All others $p(.)$ have a value of 1.
- $r(s1, \text{wait}) = r(s2, \text{wait}) = -1; r(s4, \text{wait}) = 0; r(s5, \text{wait}) = -100;$
 $r(s1, \text{move}(r1,l1,l2)) = r(s2, \text{move}(r1,l2,l1)) = -100;$
 $r(s3, \text{move}(r1,l3,l4)) = r(s4, \text{move}(r1,l4,l3)) = -100;$
 $r(s4, \text{move}(r1,l4,l5)) = r(s5, \text{move}(r1,l5,l4)) = -100;$
 $r(s1, \text{move}(r1,l1,l4)) = r(s4, \text{move}(r1,l4,l1)) = -1;$
 $r(s2, \text{move}(r1,l2,l3)) = r(s3, \text{move}(r1,l3,l2)) = -1;$
 $r(s5, \text{move}(r1,l5,l2)) = -1; r(s1) = r(s2) = r(s3) = r(s5) = 0; r(s4) = 100$

Example of an MDP

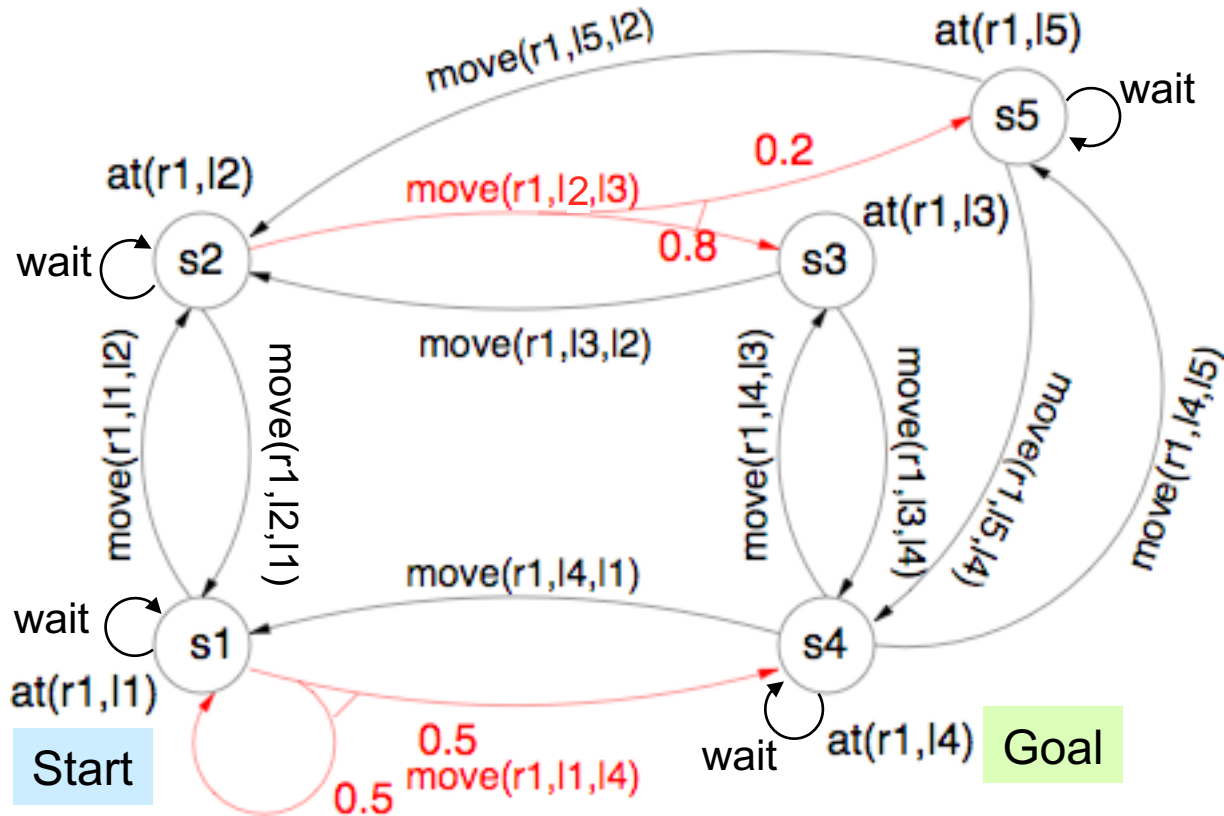


What is a Solution?

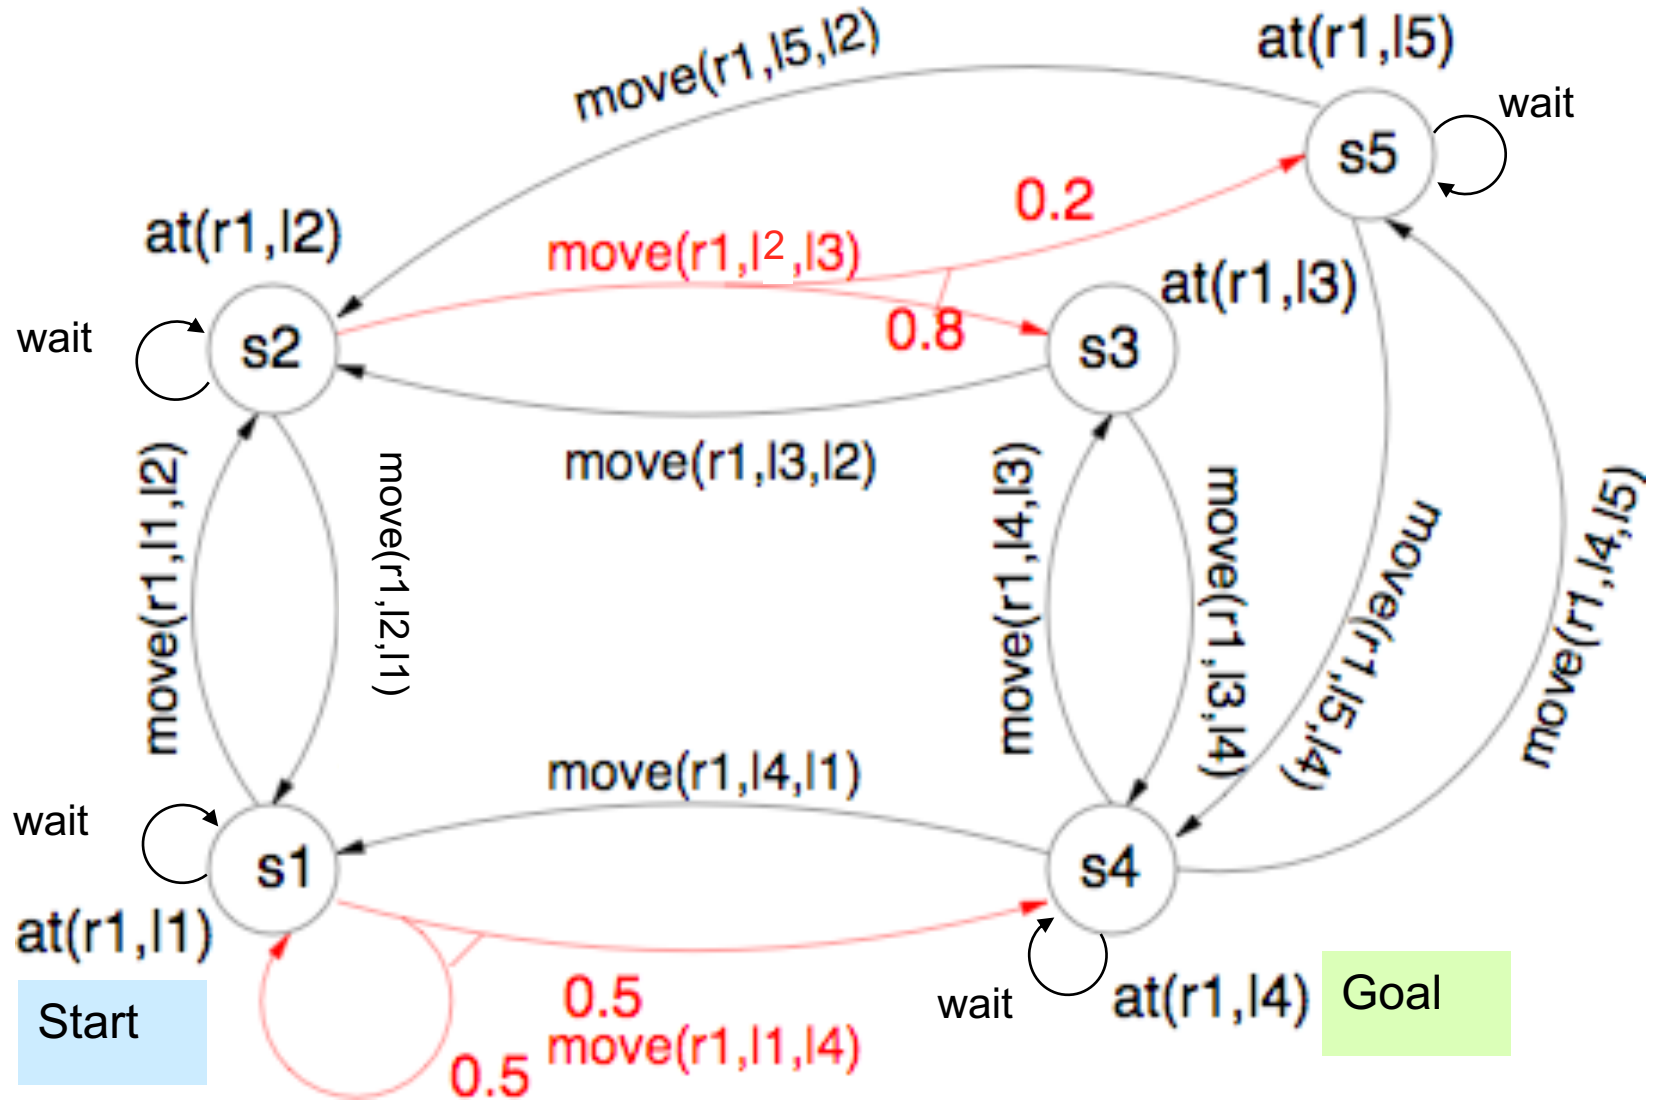
- What does a solution to the problem look like?
 - ◆ Any fixed action sequence (classical planning) will **not** solve the problem!

Example

- Robot r1 starts at location l1
 - ◆ State s1 in the diagram
- Objective is to get r1 to location l4
 - ◆ State s4 in the diagram

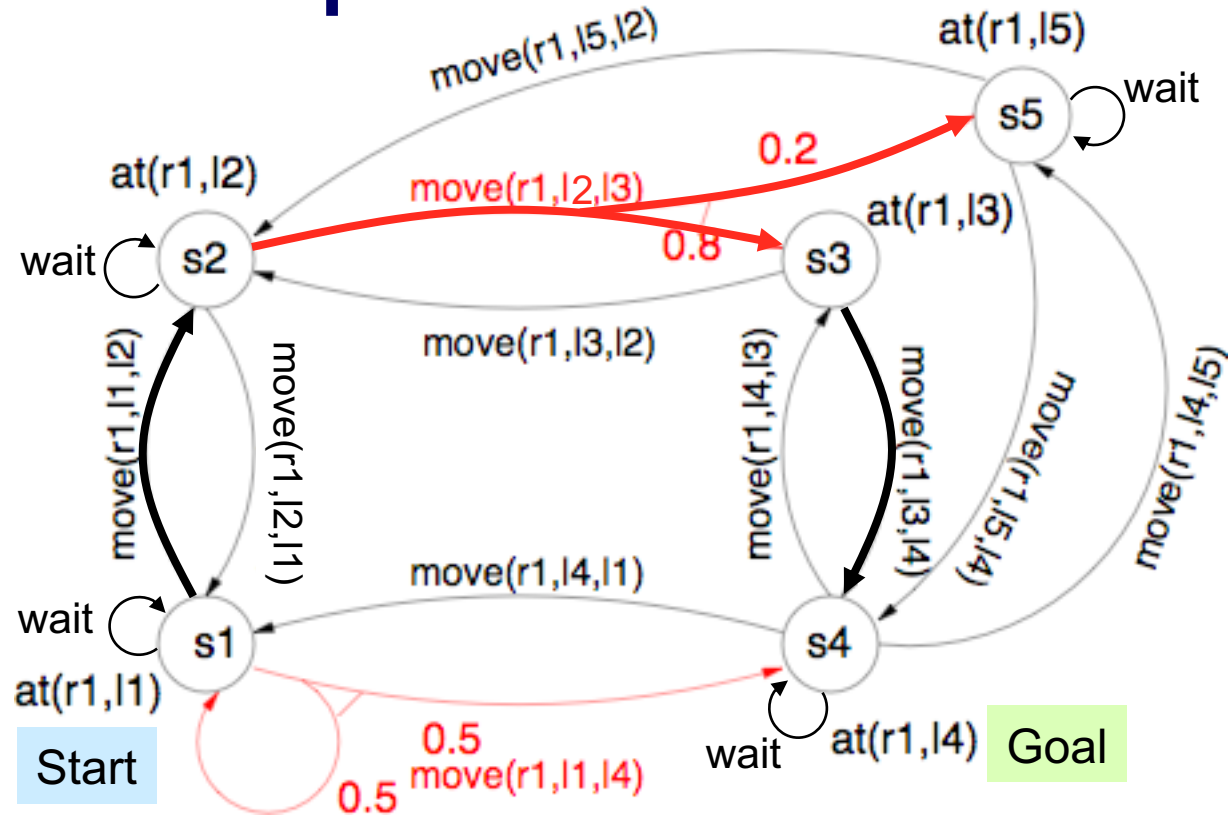


Is there a plan that will guarantee the solution?



Example

- Robot $r1$ starts at location $l1$
 - State $s1$ in the diagram
- Objective is to get $r1$ to location $l4$
 - State $s4$ in the diagram




- No fixed sequence of actions can be a solution, because we can not guarantee we will be in a state where the next action is applicable
 - e.g.,

Plan 1: $\langle \text{move}(r1, l1, l2), \text{move}(r1, l2, l3), \text{move}(r1, l3, l4) \rangle$

Example

- Robot $r1$ starts at location $l1$
 - State $s1$ in the diagram
- Objective is to get $r1$ to location $l4$
 - State $s4$ in the diagram

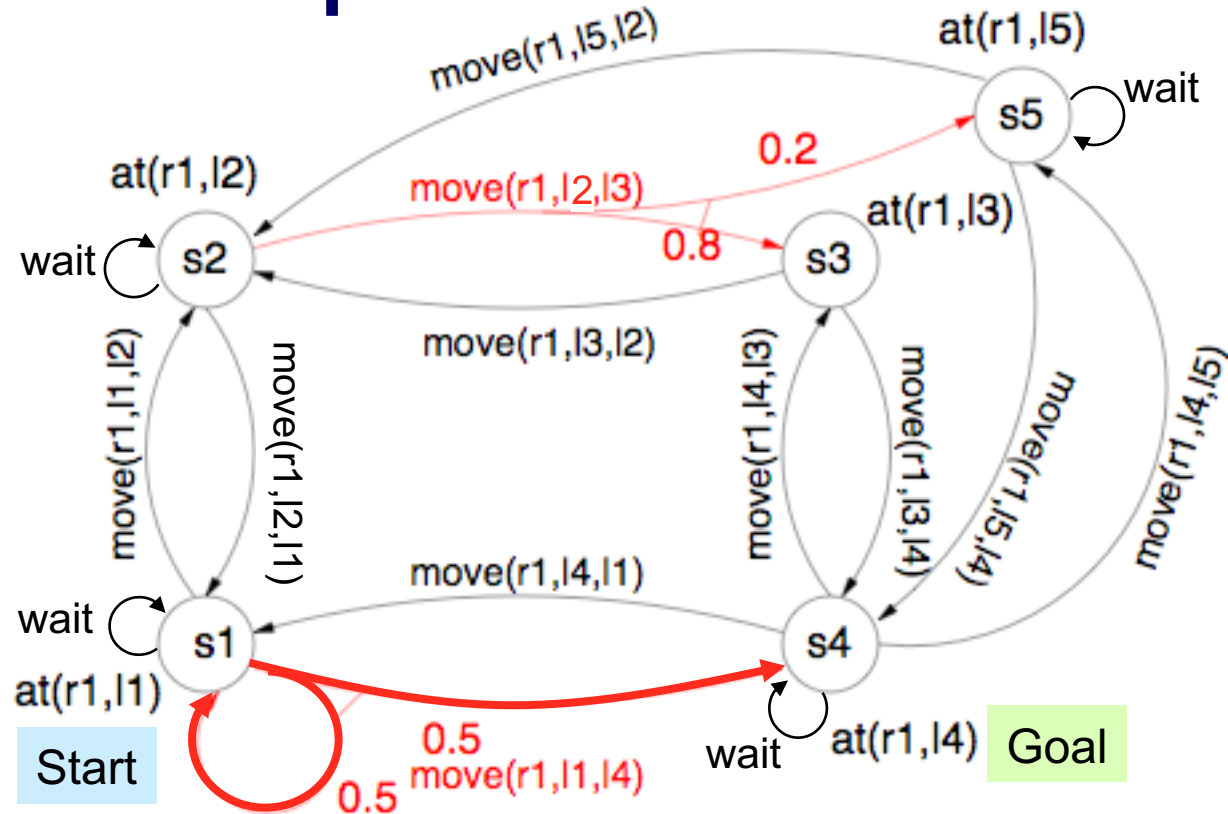


- No fixed sequence of actions can be a solution, because we can not guarantee we will be in a state where the next action is applicable
 - e.g.,

Plan 2: $\langle \text{move}(r1, l1, l2), \text{move}(r1, l2, l3), \text{move}(r1, l5, l4) \rangle$

Example

- Robot $r1$ starts at location $l1$
 - State $s1$ in the diagram
- Objective is to get $r1$ to location $l4$
 - State $s4$ in the diagram



- No fixed sequence of actions can be a solution, because we can not guarantee we will be in a state where the next action is applicable
 - e.g.,

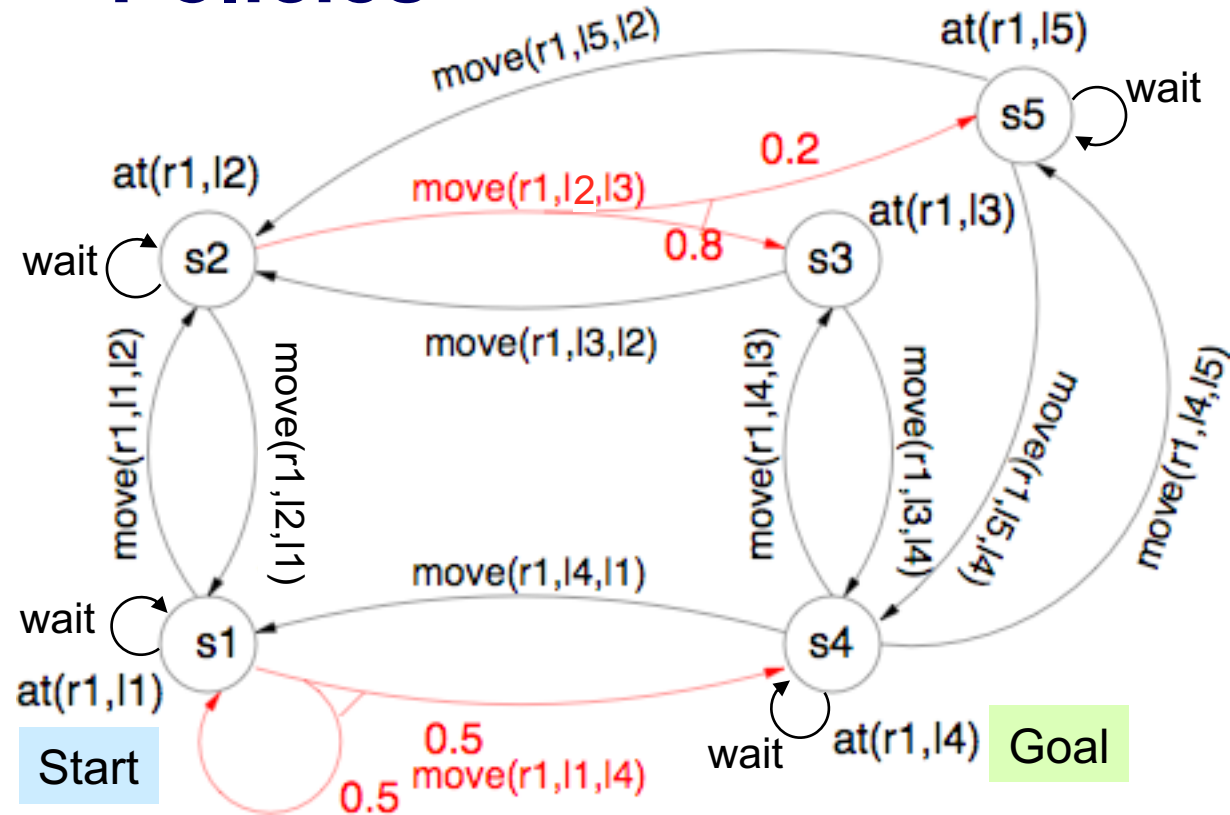
Plan 3: $\langle move(r1, l1, l4) \rangle$

What is a Solution?

- A solution must specify what the agent should do for *any* state that the agent might reach
→ *policy* π

$$\Pi: S \rightarrow A, \quad \pi(s) = a \in A$$

Policies

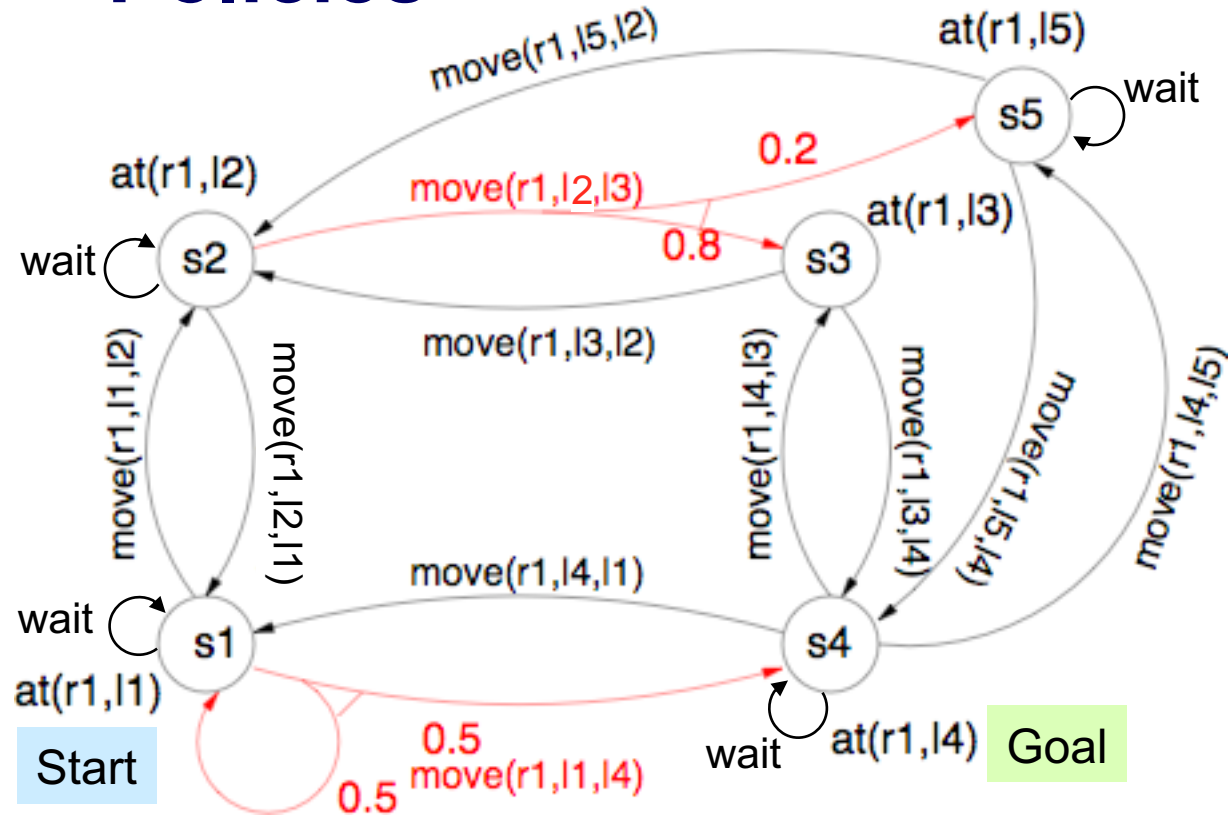


- Policy: a function that maps states into actions
- Write it as a set of state-action pairs

Policies

$$\pi_1 = \{(s1, \text{move}(r1, l1, l2)), \\ (s2, \text{move}(r1, l2, l3)), \\ (s3, \text{move}(r1, l3, l4)), \\ (s4, \text{wait}), \\ (s5, \text{wait})\}$$

$$\pi_2 = \{(s1, \text{move}(r1, l1, l2)), \\ (s2, \text{move}(r1, l2, l3)), \\ (s3, \text{move}(r1, l3, l4)), \\ (s4, \text{wait}), \\ (s5, \text{move}(r1, l5, l4))\}$$

$$\pi_3 = \{(s1, \text{move}(r1, l1, l4)), \\ (s2, \text{move}(r1, l2, l1)), \\ (s3, \text{move}(r1, l3, l4)), \\ (s4, \text{wait}), \\ (s5, \text{move}(r1, l5, l4))\}$$


- Policy: a function that maps states into actions
- Write it as a set of state-action pairs

Initial States

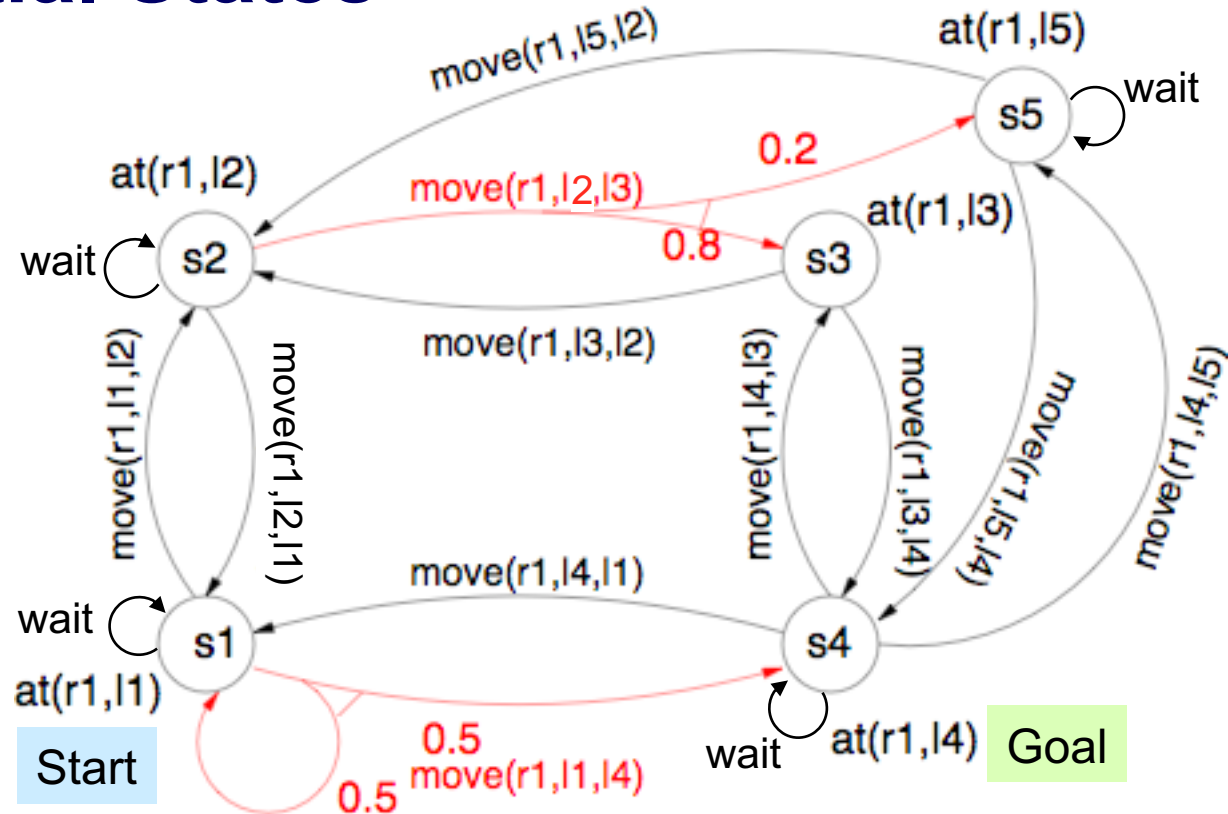
- For every state s , there will be a probability $P(s)$ that the system begins in the state s

◆ We assume the system starts in a unique initial state s_0

» $P(s_0) = 1$

» $P(s_i) = 0$ for $i \neq 0$

- In the example, $P(s_1) = 1$, and $P(s) = 0$ for all other states

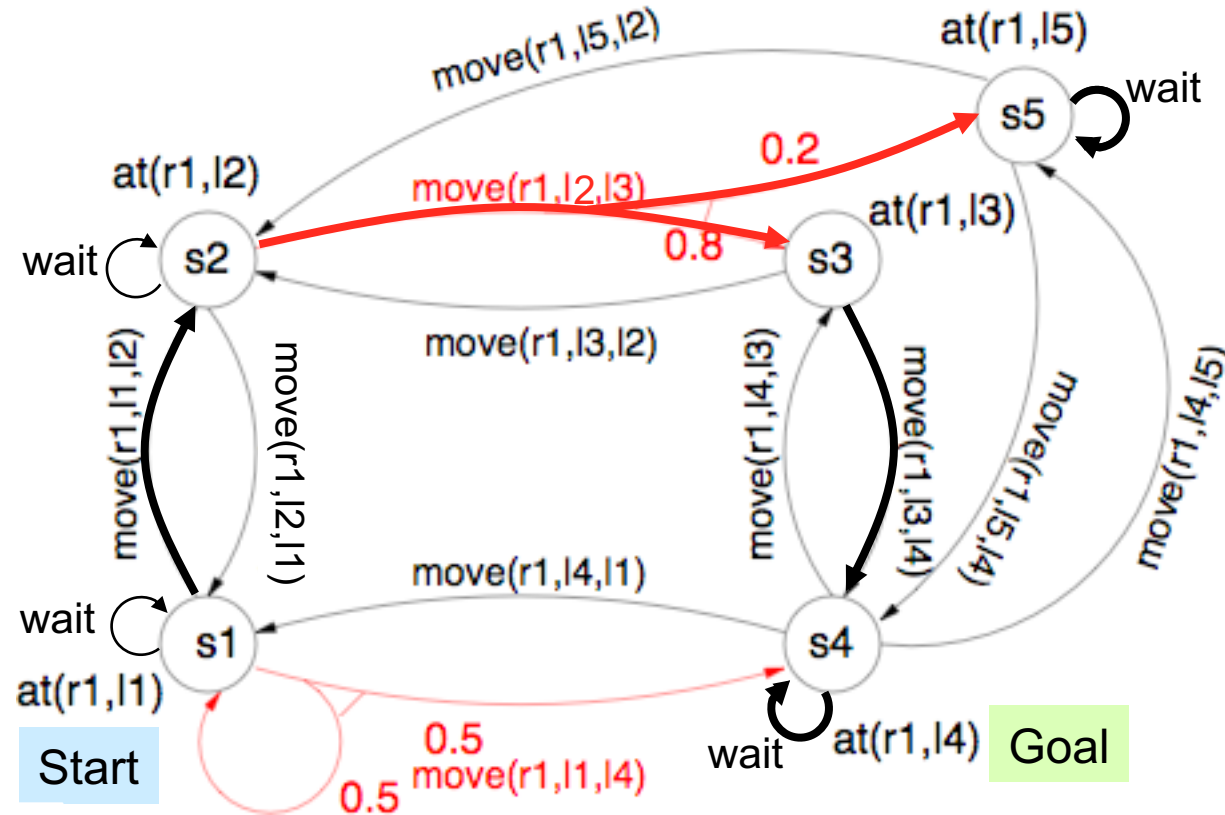


Histories

- Each time a given policy is executed starting from the initial state, the stochastic nature of the environment will lead to a different environment history.
- **Each policy induces a probability distribution over histories**
 - ◆ If $h = \langle s_0, s_1, \dots \rangle$ then
$$P(h | \pi) = P(s_0) \prod_{i \geq 0} p_{\pi(s_i)}(s_{i+1} | s_i, a_i)$$

FAZER Testinho

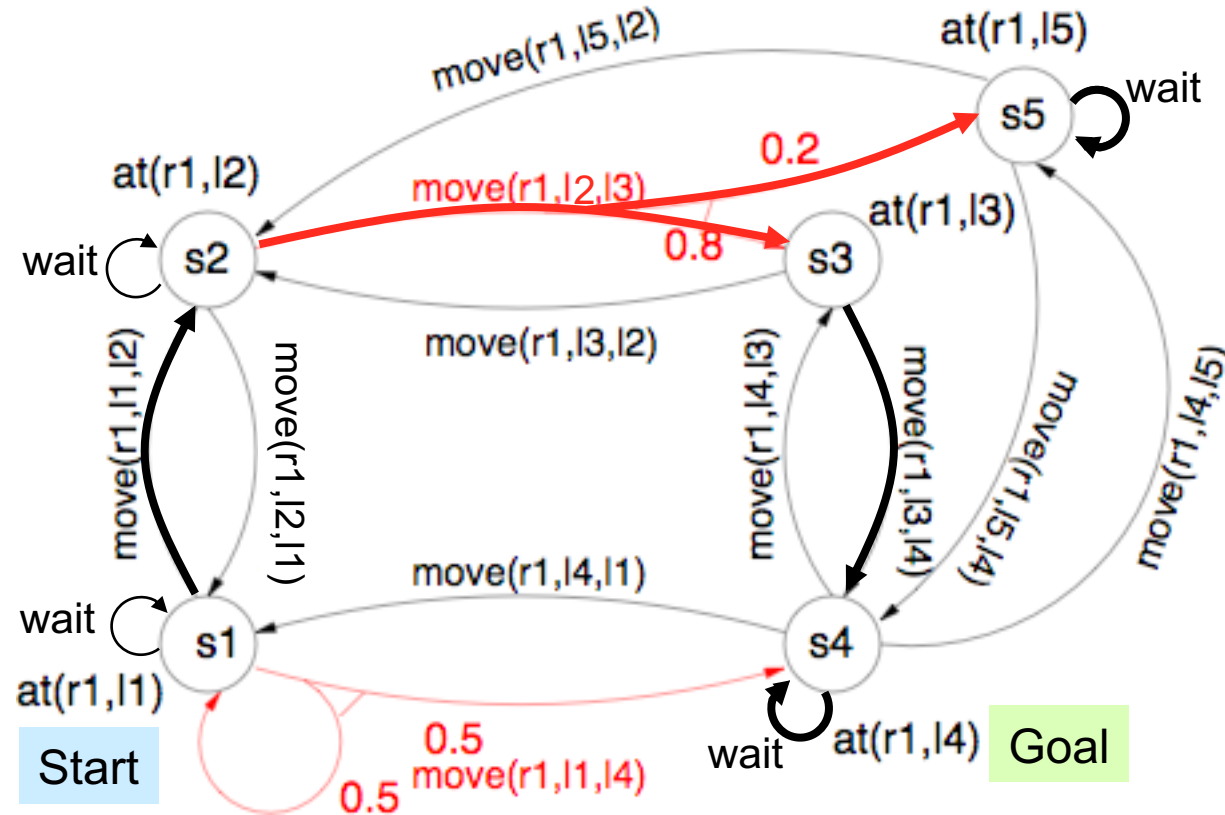
$\pi_1 = \{(s1, \text{move}(r1,l1,l2)),$
 $(s2, \text{move}(r1,l2,l3)),$
 $(s3, \text{move}(r1,l3,l4)),$
 $(s4, \text{wait}),$
 $(s5, \text{wait})\}$



Qual(is) história(s) é(são) desenvolvida(s)
 por esta política neste MDP?

FAZER TESTINHO!!!

$\pi_1 = \{(s1, \text{move}(r1, l1, l2)),$
 $(s2, \text{move}(r1, l2, l3)),$
 $(s3, \text{move}(r1, l3, l4)),$
 $(s4, \text{wait}),$
 $(s5, \text{wait})\}$



goal

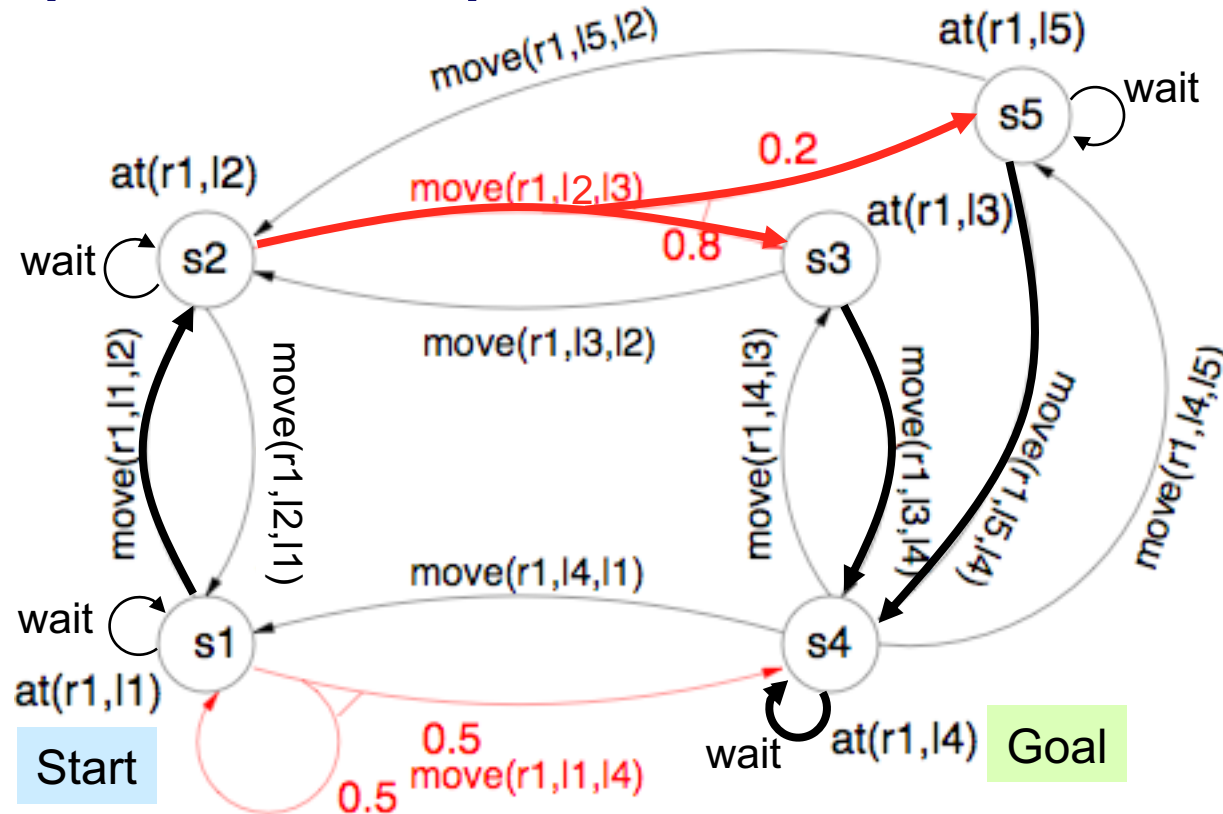
$h_1 = \langle s1, s2, s3, s4, s4, \dots \rangle$

$h_2 = \langle s1, s2, s5, s5 \dots \rangle$

Qual a probabilidade de desenvolver h_1 neste MDP?

Example (continued)

$\pi_2 = \{(s1, \text{move}(r1,l1,l2)),$
 $(s2, \text{move}(r1,l2,l3)),$
 $(s3, \text{move}(r1,l3,l4)),$
 $(s4, \text{wait}),$
 $(s5, \text{move}(r1,l5,l4))\}$



$h_1 = \langle s1, s2, s3, \text{goal}, s4, s4, \dots \rangle$
 $h_3 = \langle s1, s2, s5, \text{goal}, s4, s4, \dots \rangle$

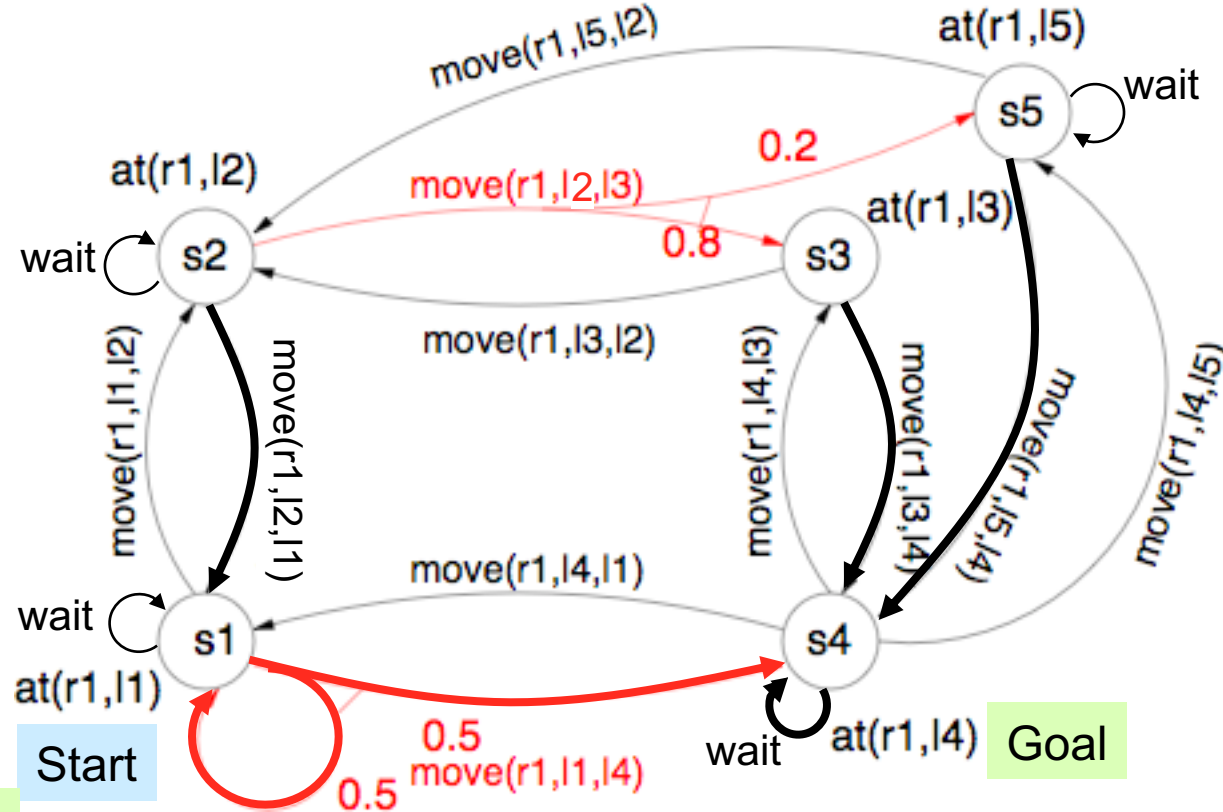
$$P(h_1 \mid \pi_2) = 1 \times 0.8 \times 1 \times \dots = 0.8$$

$$P(h_3 \mid \pi_2) = 1 \times 0.2 \times 1 \times \dots = 0.2$$

$$P(h \mid \pi_2) = 0 \text{ for all other } h$$

Example (continued)

$\pi_3 = \{(s1, \text{move}(r1, l1, l4)),$
 $(s2, \text{move}(r1, l2, l1)),$
 $(s3, \text{move}(r1, l3, l4)),$
 $(s4, \text{wait}),$
 $(s5, \text{move}(r1, l5, l4))\}$



$h_4 = \langle s1, s4, s4, \dots \rangle$

$h_5 = \langle s1, s1, s4, s4, \dots \rangle$

$h_6 = \langle s1, s1, s1, s4, s4, \dots \rangle$

...

$h_7 = \langle s1, s1, s1, s1, s1, s1, \dots \rangle$

$$P(h_4 | \pi_3) = 1 \times 0.5 \times 1 \times 1 \times 1 \times \dots = 0.5$$

$$P(h_5 | \pi_3) = 1 \times 0.5 \times 0.5 \times 1 \times 1 \times \dots = 0.25$$

$$P(h_6 | \pi_3) = 1 \times 0.5 \times 0.5 \times 0.5 \times 1 \times \dots = 0.125$$

$$P(h_7 | \pi_3) = 1 \times 0.5 \times 0.5 \times 0.5 \times 0.5 \times \dots = 0$$

Quality of a policy

- The quality of the policy is measured by the *expected utility* (or *value*) of the possible environment histories generated by that policy:

$$E[V_{\pi}(h)] = \sum_h P(h | \pi) V_{\pi}(h)$$

- An **optimal policy** π^* is a policy that yields the **highest expected utility**.


Discounted reinforcements:

$$V_{\pi}(h) = r(s_0) + \gamma r(s_1) + \gamma^2 r(s_2) + \gamma^3 r(s_3) + \dots$$

- A **discount factor** γ : $0 \leq \gamma \leq 1$

Utility Functions

- Reinforcement $r(s)$ for each state s :
 - ◆ $- : \text{cost } C(s,a)$
 - ◆ $+ : \text{reward } R(s)$
- Example:
 - ◆ $C(s,a) = 1$ for each “horizontal” action
 - ◆ $C(s,a) = 100$ for each “vertical” action
 - ◆ $C(s_1, \text{wait}) = 1; C(s_2, \text{wait}) = 1; C(s_4, \text{wait}) = 0; C(s_5, \text{wait}) = 100$
 - ◆ R as shown: $r(s_1)=r(s_2)=r(s_3)=r(s_5)=0; r(s_4)=100$
- Utility function: **generalization of a goal** (additive rewards)
 - ◆ If $h = \langle s_0, s_1, \dots \rangle$, then $V_\pi(h) = \sum_{i \geq 0} \gamma^i (R(s_i) - C(s_i, \pi(s_i)))$



Example

$\pi_1 = \{(s1, \text{move}(r1, l1, l2)),$
 $(s2, \text{move}(r1, l2, l3)),$
 $(s3, \text{move}(r1, l3, l4)),$
 $(s4, \text{wait}),$
 $(s5, \text{wait})\}$


$\gamma = 0.9$

$h_1 = \langle s1, s2, s3, s4, s4, \dots \rangle$

$$V_{\pi_1}(h_1) = .9^0(0 - 100) + .9^1(0 - 1) + .9^2(0 - 100) + .9^3 100 + .9^4 100 + \dots = 547.9$$

$h_2 = \langle s1, s2, s5, s5 \dots \rangle$

$$V_{\pi_1}(h_2) = .9^0(0 - 100) + .9^1(0 - 1) + .9^2(-100) + .9^3(-100) + \dots = -910.1$$



TESTINHO: qual a utilidade (ou valor) esperada de π_1 ?

Optimal Policy

- Utility of a state – defined in terms of the utility of state sequences:

$$V_{\pi}(s) = E \left[\sum_{t=0}^{\infty} \gamma^t r(s_t) \mid \pi, s_0 = s \right]$$

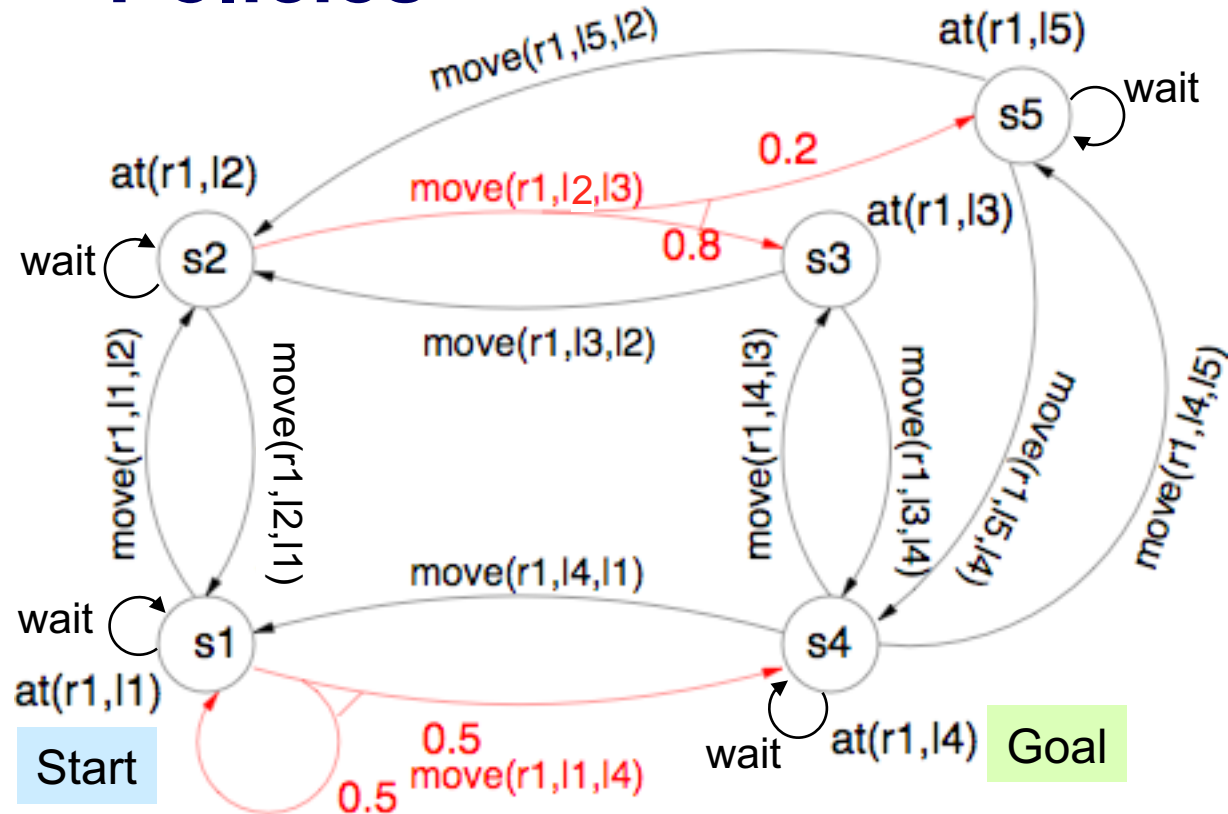
The true utility of a state, $V(s)$, is just $V_{\pi^*}(s)$, which allows the agent to choose the action that maximizes the expected utility of the *subsequent* state:

$$\pi^*(s) = \arg \max_a \sum_{s'} p(s' \mid s, a) V^*(s')$$

If we know V^* , then it's easy to find the optimal policy.

Policies

$$\pi_1 = \{(s1, \text{move}(r1, l1, l2)), \\ (s2, \text{move}(r1, l2, l3)), \\ (s3, \text{move}(r1, l3, l4)), \\ (s4, \text{wait}), \\ (s5, \text{wait})\}$$

$$\pi_2 = \{(s1, \text{move}(r1, l1, l2)), \\ (s2, \text{move}(r1, l2, l3)), \\ (s3, \text{move}(r1, l3, l4)), \\ (s4, \text{wait}), \\ (s5, \text{move}(r1, l5, l4))\}$$


Qual é melhor? π_1 ou π_2 ?

Computing V^* Approaches

- **Value iteration**
- Policy iteration
- Linear programming

Value Iteration

1. Initialize $V_0(s) = 0$, for all s .
2. Loop until a stop criterion is met:
 - ◆ Loop for all s :

$$V^{t+1}(s) \leftarrow r(s) + \max_a \gamma \sum_{s'} p(s'|s, a) V^t(s')$$

This algorithm is guaranteed to converge to V^* .
The influence of \mathbf{r} and \mathbf{p} , which we know, drives the successive V s to get closer and closer to V^* .

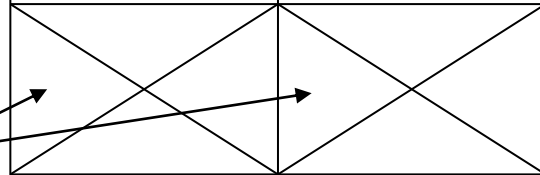
VI: Discussion

- VI computes new values in each iteration, and chooses a policy based on those values
- This algorithm converges in a polynomial number of iterations
 - ◆ But the variable in the polynomial is the number of states
 - ◆ The number of states is usually **huge**
- Need to examine the **entire state space** in each iteration
 - ◆ Thus, this algorithm takes huge amounts of time and space

TAREFA

- Ambiente discreto 4x4, com obstáculos.
- Agente deve alcançar posição destino **D** a partir de **qualquer lugar** do ambiente.
- **D** é um estado absorvente (ao atingir **D**, o episódio termina): $V^*(D) = 0$
- Ações que o agente pode realizar: N, S, L, O
- Penalidade por executar uma **ação** (qualquer) = -1
 - ◆ Melhor política \Rightarrow caminho mais curto
- Considerar $\gamma = 1$ e MDP determinístico ($p=1$)

Ambiente

(1,4)	(2,4)	(3,4)	(4,4)
(1,3)	(2,3)	(3,3)	(4,3)
		(3,2)	(4,2)
		(3,1)	(4,1)

D

Ações = {N, S, L, O}

Tarefa: algoritmo de iteração de valor para MDP determinístico

- Cálculo iterativo da função valor ótima.

$$V(s) \leftarrow r_{s,a} + \max_a (V(s'))$$

Repetir até $V(s)$ estabilizar.

Sendo:

s – estado atual, s' – próximo estado,

$r_{s,a}$ – reforço recebido por executar a em s

$V(.)$ – valor do estado

Exemplo de cálculo de $V(s)$

Início

0	0	s'_2 0	0
0	s'_1 0	s 0	s'_3 0
		s'_4 0	0
D 0	0	0	0

Iteração 1 (quando calculous para todos os estados)

		-1	
D 0			

■ ■ ■

$$V(s) = \max_a ((r(s, O) + V(s'_1)), \\ (r(s, N) + V(s'_2)), \\ (r(s, L) + V(s'_3)), \\ (r(s, S) + V(s'_4)))$$

$$= \max_a ((-1+0), (-1+0), (-1+0), (-1+0)) \\ = -1$$

Tarefa

- Entrega: Mostrar o valor no espaço de estados (grade com o valor em cada célula) após CADA ITERAÇÃO, até a convergência do algoritmo VI
- Responder:
 - ◆ Qual estado tem valor **mínimo**? Qual o valor deste estado?
 - ◆ Qual estado tem valor **máximo**? Qual o valor deste estado?
 - ◆ Mostrar na grade qual é a **política ótima** em cada célula.